

# Executing code on columnar data: the translation problem and formats that help



Jim Pivarski

Princeton – DIANA

February 8, 2017

I'm working on a query language and database server to aggregate large samples of HEP data on the fly.

**Purpose:** to eliminate the need for private skims in most situations.



The query language, Femtocode, plays a similar role as TTreeFormula:

- ▶ a high-level language for the physicist
- ▶ usually for filling a histogram (so query responses are small)
- ▶ but generally useful for transforming one dataset into another.

However, it's a full-fledged language with assignments and user-defined functions, so that it can encompass a larger part of the analysis.

(I've examined SQL, LINQ, and others, and they are not sufficient. I would use a standard if I could.)

The essential feature of Fentocode is that it can transform complex structure-manipulations, which would ordinarily have to be performed in object-oriented code, into a series of vectorized kernels.

It operates on columns.

## Example:

```
hist = dataset..bin(100, 0, 50, ""  
    muons.map(m => sqrt(m.px**2 + m.py**2)).max()  
    """)
```

## transforms to

1. Pass over all muons, computing  $\sqrt{p_x^2 + p_y^2}$  on all muons, ignoring event boundaries.
2. Find the maximum such value for each event.
3. Bin those events and fill the histogram.

## rather than

1. Loop over events:
  - 1.1 Loop over muons:
    - 1.1.1 Compute  $\sqrt{p_x^2 + p_y^2}$  for each.
  - 1.2 Fill a histogram with the maximum.

## Three types of data transformations:

**Flat:** apply  $N$ -argument function to each element of  $N$  aligned arrays.

Known in the Numpy community as a “ufunc.”

**Explode:** emulate (nested) for loops by replicating data in one array so that it becomes aligned with another array.

**Reduce:** emulate reducer functions (sum, mean, max. . . ) by combining elements of an array so that it becomes aligned with an outer level of structure.

## Three types of data transformations:

**Flat:** apply  $N$ -argument function to each element of  $N$  aligned arrays.

Known in the Numpy community as a “ufunc.”

**Explode:** emulate (nested) for loops by replicating data in one array so that it becomes aligned with another array.

**Reduce:** emulate reducer functions (sum, mean, max. . . ) by combining elements of an array so that it becomes aligned with an outer level of structure.

## What's missing?

“Repeat until convergence.” Whatever determines “convergence” may be different for each element of the array: they’d all wait for the last one, anyway.

The majority of steps in a typical calculation are flat:

```
double in[ZILLION];  
double out[ZILLION];  
  
for (int i = 0; i < ZILLION; i++)  
    out[i] = flat_operation(in[i]);
```

- ▶ Compilation with `-O3` does auto-vectorization if possible (depends on `flat_operation`).
- ▶ Easiest form for CPU to prefetch memory and/or pipeline operations.
- ▶ This form is also ideal for GPU calculations.
- ▶ There is a standard for functions like this: Numpy's `ufunc` is widely used among scientific libraries.
  - ▶ Easy way for a user to add functions to the language!



```
import ctypes, numpy, numba

libMathCore = ctypes.cdll.LoadLibrary("libMathCore.so")
chi2_ctypes = libMathCore._ZN5TMathl7ChisquareQuantileEdd # c++filt!
chi2_ctypes.argtypes = (ctypes.c_double, ctypes.c_double)
chi2_ctypes.restype = ctypes.c_double

# compile to pure-C ufunc
@numba.vectorize(["f8(f8, f8)"], nopython=True)
def chi2_ufunc(p, ndf):
    return chi2_ctypes(p, ndf)

p = numpy.random.uniform(0, 1, int(1e6)) # million random numbers
result = chi2_ufunc(p, 100) # call ufunc on all of them
# 3.22 seconds

import ROOT
result = [ROOT.TMath.ChisquareQuantile(pi, 100) for pi in p]
# 9.32 seconds
```

(Performance comparison is just to show that the ufunc computes ChisquareQuantile in C, not in Python. Simpler functions show a more dramatic difference.)