

Executing code on columnar data: the translation problem and formats that help

Jim Pivarski

Princeton – DIANA

February 8, 2017

Three types of data transformations:

Flat: apply N -argument function to each element of N aligned arrays.

Known in the Numpy community as a “ufunc.”

Explode: emulate (nested) for loops by replicating data in one array so that it becomes aligned with another array.

Reduce: emulate reducer functions (sum, mean, max. . .) by combining elements of an array so that it becomes aligned with an outer level of structure.

Three types of data transformations:

Flat: apply N -argument function to each element of N aligned arrays.

Known in the Numpy community as a “ufunc.”

Explode: emulate (nested) for loops by replicating data in one array so that it becomes aligned with another array.

Reduce: emulate reducer functions (sum, mean, max. . .) by combining elements of an array so that it becomes aligned with an outer level of structure.

What's missing?

“Repeat until convergence.” Whatever determines “convergence” may be different for each element of the array: they’d all wait for the last one, anyway.

The majority of steps in a typical calculation are flat:

```
double in[ZILLION];  
double out[ZILLION];  
  
for (int i = 0; i < ZILLION; i++)  
    out[i] = flat_operation(in[i]);
```

- ▶ Compilation with `-O3` does auto-vectorization if possible (depends on `flat_operation`).
- ▶ Easiest form for CPU to prefetch memory and/or pipeline operations.
- ▶ This form is also ideal for GPU calculations.
- ▶ There is a standard for functions like this: Numpy's `ufunc` is widely used among scientific libraries.
 - ▶ Easy way for a user to add functions to the language!

```
import ctypes, numpy, numba

libMathCore = ctypes.cdll.LoadLibrary("libMathCore.so")
chi2_ctypes = libMathCore._ZN5TMathl7ChisquareQuantileEdd # c++filt!
chi2_ctypes.argtypes = (ctypes.c_double, ctypes.c_double)
chi2_ctypes.restype = ctypes.c_double

# compile to pure-C ufunc
@numba.vectorize(["f8(f8, f8)"], nopython=True)
def chi2_ufunc(p, ndf):
    return chi2_ctypes(p, ndf)

p = numpy.random.uniform(0, 1, int(1e6)) # million random numbers
result = chi2_ufunc(p, 100) # call ufunc on all of them
# 3.22 seconds

import ROOT
result = [ROOT.TMath.ChisquareQuantile(pi, 100) for pi in p]
# 9.32 seconds
```

(Performance comparison is just to show that the ufunc computes ChisquareQuantile in C, not in Python. Simpler functions show a more dramatic difference.)